

Project 1 - Matrix Class & Operator Overloading

CSCI 1730 – Fall 2016

DUE MON 2016-09-12 @ 11:55 PM

Problem / Exercise

Your goal for this breakout is to create an easy-to-use **Matrix** class in C++ that makes use of dynamic memory allocation and includes basic matrix operations both in the form of regular functions and via operator overloading. Certain aspects of this project may not be covered in lecture. You are encouraged to use the DEITEL text and cpreference.com as references. Also, please ask questions via Piazza. Here is the basic prototype for the **Matrix** class (you may need to add more to it to support some of the additional features listed further down in this document):

```
typedef unsigned int uint;
typedef initializer_list<initializer_list<double>> i_list;

class Matrix {

public:

    Matrix(uint rows, uint cols);           // constructor (all elements initialized to 0)
    Matrix(const i_list & list);            // constructor (using initializer list)
    Matrix(const Matrix & m);              // copy constructor
    ~Matrix();                             // destructor

    Matrix add(double s) const;             // add scalar to this matrix
    Matrix add(const Matrix & m) const;     // add this matrix and another matrix

    Matrix subtract(double s) const;        // subtract scalar from this matrix
    Matrix subtract(const Matrix & m) const; // subtract another matrix from this matrix

    Matrix multiply(double s) const;        // multiply this matrix by a scalar
    Matrix multiply(const Matrix & m) const; // multiply this matrix by another matrix

    Matrix divide(double s) const           // divide this matrix by a scalar
    Matrix t() const;                      // transpose of this matrix

    const uint numRows() const;            // returns the number of rows
    const uint numCols() const;           // returns the number of cols

    double & at(uint row, uint col);       // get/set element at row,col
    const double & at (uint row, uint col) const; // get element at row,col (when using a const object)

}; // Matrix
```

Important Class Details

Your **Matrix** implementation will contain elements of type **double**. In the prototype presented above, the term scalar refers to a regular number. For example, if you add a scalar to a matrix, then each element in the matrix gets that number added to it. Contrast this with the member functions that take a **Matrix** as their parameter. These functions represent regular matrix operations. For some of these operations (e.g., multiplication, transpose, etc.), you may need to consult some sort of reference in order to recall the exact procedure/meaning behind the operation.

NOTE: You **MAY** assume valid input for all operations.

NOTE: You **MAY NOT** use library classes such as `std::array`, `std::vector`, `std::list`, etc. for this project. You must implement your `Matrix` class internally using a dynamically allocated array.

Example Usage 1

```
Matrix a(2, 2);
a.at(0, 0) = 1; // [ 1, 2 ]
a.at(0, 1) = 2; // [ 1, 3 ]
a.at(1, 0) = 1;
a.at(1, 1) = 3;

Matrix b(2, 1);
a.at(0, 0) = 3; // [ 3 ]
a.at(0, 1) = 2; // [ 2 ]

Matrix c = a.multiply(b);
cout << "[ " << c.at(0, 0) << " ]" << endl // [ 7 ]
      << "[ " << c.at(1, 0) << " ]" << endl; // [ 9 ]

Matrix d = {{1, 2}, // this will implicitly call the overloaded constructor
            {3, 4}}; // that takes an initializer list
```

The usage of the member function `at(uint, uint)` is what facilitates our ability to perform operations such as `a.at(0, 0) = 1`. If you implement this function carefully, then this behavior should work because the function returns a reference to an element. In order to support the constructor overload (i.e., matrix construction using an initializer list), you will need to use a standard template library (STL) class called `std::initializer_list`¹. The type signature for the `<<` parameter representing the list should be `std::initializer_list<std::initializer_list<double>>` or simply `i_list` if you use the provided `typedef`

Operator Overloading

You will also need to overload operators in order to support the following functionality. It is up to you whether or not these should be member or non-member overloads.

```
// assume we have two matrices of appropriate size already set up
Matrix a;
Matrix b;

// after providing the overloads, you should be able to do any of the following operations
// using regular operators instead of the member functions
Matrix c0 = a + 5.2;
Matrix c1 = a + a; // NOTE: these examples actually end up calling the copy constructor
Matrix c2 = a - 3.5;
Matrix c3 = b - b;
Matrix c4 = a * 2.1;
Matrix c5 = a * b;
Matrix c6 = a / 2.0;
```

¹`std::initializer_list`: http://en.cppreference.com/w/cpp/utility/initializer_list

You should also support stream insertion (similar to overriding the `toString` method in Java) so that your matrices can easily be printed.

```
cout << a << endl; // example output: [ 1, 2 ]
                //                  [ 1, 3 ]
```

You should also support matrix assignment using an initializer list (to easily overwrite existing elements) that looks like the following:

```
Matrix d(2, 2);
d = {{ 1, 2 },
     { 3, 4 }};
```

In order to support this last operator overload (i.e., matrix assignment using an initializer list), you will need to use a standard template library (STL) class called `std::initializer_list`. The type signature for the `<<` parameter representing the list should be `std::initializer_list<std::initializer_list<double>>` or simply `i_list` if you use the provided `typedef`. It is preferred that you specify the parameter as a `const i_list &` in order to avoid any unnecessary copying.

Additional Features

In addition to the requirements for listed above, you need to make sure your `Matrix` class supports the following features:

- **Overloaded Function Call Operator (`operator()(uint row, uint col)`):** After creating a (non-dynamically allocated) `Matrix` object, the user should be able to access the elements using the function call operator (as an alternative to using the `at` function):

```
Matrix a(1, 1);
a(0, 0) = 5;
cout << a(0,0) << endl;
```

- **Overloaded Copy Assignment Operator (`operator=(const Matrix &)`):** You should have already overloaded the assignment operator to take in a special kind of initializer list. Now you need to provide an additional overload that supports copy assignment. This will make your `Matrix` class more consistent since copy assignment parallels copy construction. Here is an example:

```
Matrix a(1, 1);
a(0, 0) = 5;

Matrix b(1, 1);
b = a; // copy assignment
```

- **Overloaded Non-Member Arithmetic Operators for Scalars:** You should have already created overloads to support the basic arithmetic operations where the right-hand-side of an operation is a scalar value. Now you need to implement operator overloads so that scalars can be used on the left-hand-side of an operation. Here is an example showing the operators that you need to support:

```
Matrix a = {{1, 2},
            {3, 4}};

Matrix b = 4.0 + a; // [ 5, 6 ]
                // [ 7, 8 ]

Matrix c = 4.0 - a; // [ 3, 2 ]
                // [ 1, 0 ]

Matrix d = 2.0 * a; // [ 2, 4 ]
                // [ 6, 8 ]

Matrix e = 12.0 / a; // [ 12, 6 ]
                [ 4, 3 ]
```

- **Overloaded Unary Minus Operator (`operator-()`):** You need to support negating your `Matrix` objects:

```
Matrix a = {{1, 2}};
cout << -a << endl; // [ -1, -2 ]
```

1 C++ Code & Program

1.1 Setup

Make sure that all of your files are in a directory called `LastName-FirstName-p1`, where `LastName` and `FirstName` are replaced with your actual last and first names, respectively.

1.2 Source Code Files

You should organize your project into the following files:

- **Matrix.h**: This file should include the class prototype presented above as well as the prototypes for operator overloads that you implement. You **MAY NOT** modify the function prototypes that are included in the **Matrix** class prototype. However, you may add additional function prototypes and variables to the class prototype as needed. Make sure that this header file also includes a header guard (i.e., the `#ifndef` macro, etc.).
- **Matrix.cpp**: This file should contain the implementation of your class's functions as well as the implementation of any operator overloads that you implement.
- **p1.cpp**: This file should contain a small/moderately sized driver that demonstrates the full range of functionality of your **Matrix** class.

Additionally, make sure that you adhere to the following:

- All functions must be documented using Javadoc-style comments. Use inline documentation, as needed, to explain ambiguous or tricky parts of your code.

1.3 Makefile File

You need to include a **Makefile**. Your **Makefile** needs to compile and link separately. Make sure that your **Matrix.cpp** file compiles to **Matrix.o**. This is very important because we will be testing your submission by linking against your **Matrix.o** file. The resulting executable should be called **p1**.

Make sure that when you compile, you pass the following options to **g++** in addition to the **-c** option:

```
-Wall -std=c++14 -g -O0 -pedantic-errors
```

Here is a link to the **Makefile** for the "gradebook" example we coded up in class that you might find useful as a reference point: <https://gist.github.com/mepcotterell/45a10fa72b208a76d968>.

1.4 README File

Make sure to include a **README** file that includes the following information presented in a reasonably formatted way:

- Your Name and 810/811#
- Instructions on how to compile and run your program.

Here is a partially filled out, example **README** file: <https://gist.github.com/mepcotterell/3ce865e3a151a3b49ec3>.

NOTE: Try to make sure that each line in your **README** file does not exceed 80 characters. Do not assume line-wrapping. Please manually insert a line break if a line exceeds 80 characters.

1.5 Compiler Warnings

Since you should be compiling with both the **-Wall** and **pedantic-error** options, your code is expected to compile without **g++** issuing any warnings. For this project, compiling without warnings will be one or more of the test cases.

1.6 Memory Leaks

Since this project makes use of dynamic memory allocation, you are expected to ensure that your **Matrix** implementation doesn't result in any memory leaks. We will test for memory leaks using the **valgrind** utility. For this project, having no memory leaks will be one or more of the test cases.

2 Submission

Make sure your work is on `nike.cs.uga.edu` in a directory called `LastName-FirstName-p1`. From within the parent directory, execute the following command:

```
$ submit LastName-FirstName-p1 cs1730a
```

It is also a good idea to email a copy to yourself. To do this, simply execute the following command, replacing the email address with your email address:

```
$ tar zcvf LastName-FirstName-p1.tar.gz LastName-FirstName-p1
$ mutt -s "p1" -a LastName-FirstName-p1.tar.gz -- your@email.com < /dev/null
```